

Birdhouse - supporting Web Processing Services

Carsten Ehbrecht¹ Stephan Kindermann¹ Nils Hempelmann²

¹DKRZ - German Climate Compute Center

²GIZ - German Development Cooperation

28th of November 2017/ Python Frameworks Workshop at ECMWF

Motivation

GET

Post

SOAP




GetCapabilities

DescribeProcess

Execute

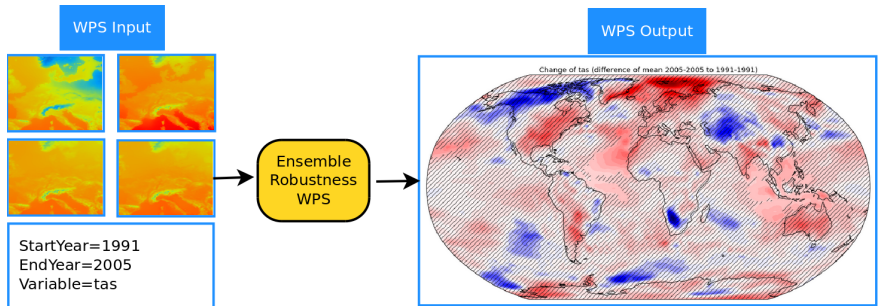
[http://www.slideshare.net/TheodorFoerster/
restful-web-processing-service](http://www.slideshare.net/TheodorFoerster/restful-web-processing-service)

WPS Use Case



WPS for Point Clouds by Adam Steer, NCI, Australia

WPS Inputs and Outputs



```
@wps # wps decorator
def myplot(nc_file , variable):
    """
    Generates a plot for given dataset and variable.

    nc_file application/netcdf Dataset
    variable string Variable name
    """
    # ... create a nice plot here
    return plot.png
```

Running Function as Web Processing Service

```
$ curl -s -o result.xml \  
http://localhost/wps? \  
  &service=WPS \  
  &version=1.0.0 \  
  &request=execute \  
  &identifier=myplot \  
  &DataInputs=nc_file=http://;variable=tas
```



PyWPS

What is PyWPS?



- An implementation of the OGC Web Processing Service standard
- Implements WPS 1.0.0 standard (WPS 2.0.0 in progress)
- Coded in the Python language (researcher friendly)
- Easy to hack (developer friendly)
- Relevant contributions by over a dozen individuals
- OSGeo accreditation around the corner ...

<http://pywps.org>



- Nginx - security, reverse proxy, load balancing
- Gunicorn - concurrency, WSGI
- WSGI App - PyWPS WSGI application

It's complicated!

- Scalability requires the orchestration of various software layers:
 - Nginx
 - Gunicorn
 - PyWPS
- Many packages to install
- Many configurations to set-up
- No clients for testing

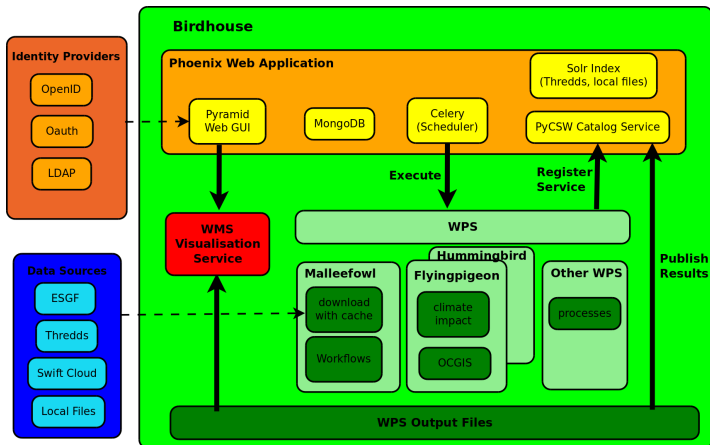
Too much work?

Birdhouse

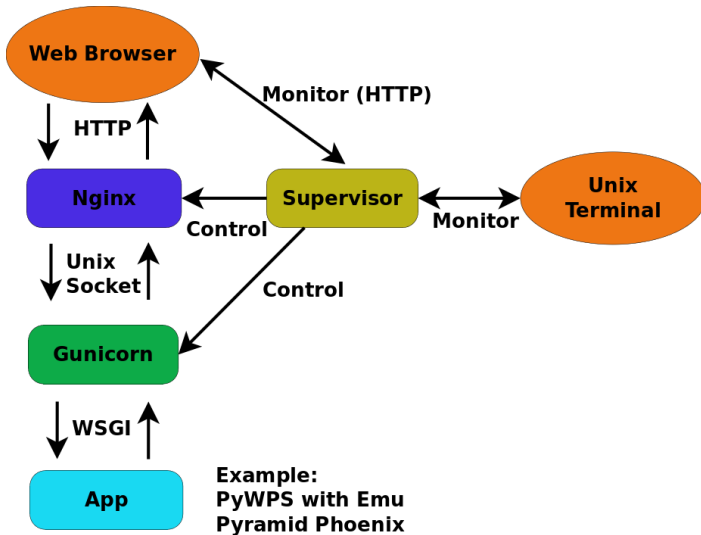
What is Birdhouse?

- Supporting OGC Web Processing Services in the climate science community.
- Making it easier to set-up WPS services (Birdhouse-Builder).
- Providing Python library and WPS processes to access climate data.
- Providing a security proxy middleware to protect OGC/WPS services.
- Providing a web application and command-line client to interact with WPS services.
- `http://bird-house.github.io/`

Birdhouse Overview



WSGI Application controlled by Supervisor





- Many components: PyWPS, Nginx, Gunicorn, ncWMS, ...
- Lots of packages: cdo, cfchecker, OCGIS, numpy, R, ESMValTool, ...
- Many configurations to set-up
- Reproducible installation
- Different Linux distributions (Centos, Debian, ...)

Conda Package Manager

- Originally for python ... but has a general concept
- Does not need admin rights

Install PyWPS from birdhouse channel

```
$ conda install -c birdhouse pywps
```

Create conda environment **pywps**

```
$ conda create -n pywps -c birdhouse pywps
```

Deployment with Buildout

- Makefile to wrap Buildout commands
- A common deployment pattern for all Birds

```
$ git clone https://github.com/bird-house/emu.git
$ cd emu
$ make clean install
$ make start
```

Update configuration like hostname, port

```
$ vim custom.cfg
$ make update
$ make restart
```

- Docker Hub is a public repository for dockerfiles
- Birdhouse repository with automatically updated docker images

`https://hub.docker.com/u/birdhouse/`

Running a docker image with Emu WPS on port 8080



```
$ docker pull birdhouse/emu
$ docker run -it -p 8000:8000 -p 8080:8080 \
    --name=emu_wps birdhouse/emu
```

Running a WPS GetCapabilities request

```
$ curl -s -o caps.xml \
http://localhost:8080/wps? \
    service=WPS& \
    version=1.0.0& \
    request=GetCapabilities
```

Phoenix web-based WPS client

PHOENIXProcessesHelp


Sign In


Phoenix


A Python Pyramid Web Application
to interact with Web Processing Services


Highlighted Processes


Run one of these favorite processes or explore [more](#).



sleep


ncdump


hello



wordcounter


cchecker


spotchecker

Explore Phoenix

Making it easy to run processes from a Web Processing Service and to visualize and share the results.



Phoenix Example


Choose Cloud Taylor

[illegible]

Submit Cloud Taylor

[illegible]

Output: Cloud Taylor Diagram



Show
Output

Monitor
Job
Execution

Job Running [20/100]

```
1 0:00:03 04c PyWPS Process cloud_taylor accepted
2 0:00:05 04c running done
```

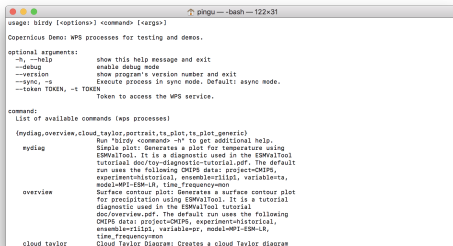
Job
Succeeded

Job Succeeded

```
1 0:00:02 %n PyDP Process cloud_taylor accepted
2 0:00:05 %n PyDP Process cloud_taylor accepted
3 0:00:07 %n PyDP Process cloud_taylor accepted
4 0:00:09 %n PyDP Process cloud_taylor accepted
5 0:00:11 %n PyDP Process cloud_taylor accepted
6 0:00:16 %n PyDP Process cloud_taylor accepted
7 0:00:21 %n PyDP Process cloud_taylor accepted
8 0:00:26 %n PyDP Process cloud_taylor accepted
9 0:00:26 100% PyDP Process Cloud Taylor Stream finished
```

Birdy command line WPS client

```
$ conda install -c birdhouse birdhouse-birdy
$ export WPS_SERVICS=http://localhost:8094/wps
$ birdy -h
```



```
usage: birdy [<options>] <command> [<args>]

Copernicus Demo: WPS processes for testing and demos.


optional arguments:
  -h, --help            show this help message and exit
  --debug               enable debug mode
  --version             show program's version number and exit
  --sync, -s           Execute process in sync mode. Default: asyene mode.
  --token TOKEN, -t TOKEN
                        Token to access the WPS service.

command:
  List of available commands (wps processes)

{mydiag,overview,cloud_taylor,portrait,ts_plot,ts_plot_generic}
  Run "birdy <command> -h" to get additional help.
  esdiag               Simple plot: Generates a plot for temperature using
                        ESMValTool. It is a diagnostic used in the ESMValTool
                        tutorial doc/toy-diagnostic-tutorial.pdf. The default
                        run uses the following CMIP5 data: project=CMIP5,
                        experiment=historical, ensemble=r1i1p1, variable=ts,
                        model=WPI-ESM-LR, time_frequency=mon
  overview             Surface contour plot: Generates a surface contour plot
                        for precipitation using ESMValTool. It is a tutorial
                        diagnostic used in the ESMValTool tutorial
                        doc/overview.pdf. The default run uses the following
                        CMIP5 data: project=CMIP5, experiment=historical,
                        ensemble=r1i1p1, variable=pr, model=WPI-ESM-LR,
                        time_frequency=mon
  cloud_taylor         Cloud Taylor Diagram: Creates a cloud Taylor diagram
```

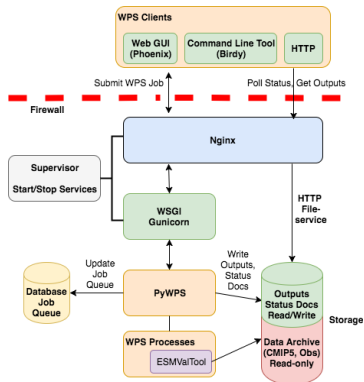

WPS Security Proxy

- Using string token (uuid) as part of URL or in request header to protect WPS execute access
- X509 certificates to access (remote) data from ESGF are provided by proxy (using environ)
- Implemented as WSGI middleware service
- `http://twitcher.readthedocs.io/en/latest/`



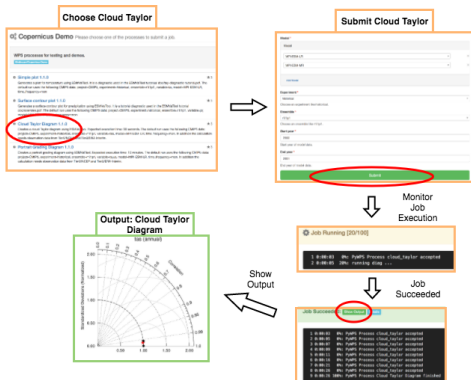
Examples

Copernicus: ESMValTool diagnostics as WPS



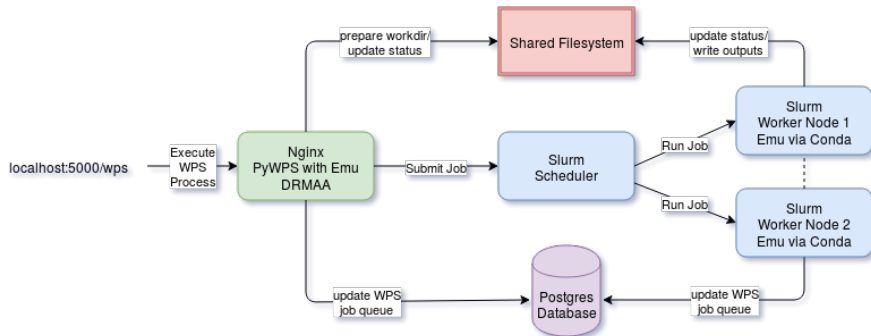
<https://github.com/cp4cds/copernicus-wps-demo>

Copernicus: Example Run



<https://bovec.dkrz.de/processes/list?wps=copernicus>

Copernicus: PyWPS Scheduler Extension



This extension is part of PyWPS

PAVICS: A Platform for the Analysis and Visualization of Climate Science

- Project based on Birdhouse by Ouranos and CRIM, Canada
- Ouranos needs a platform for climate services
- Creating and delivering climate products
- Make climate research less painful

<https://www.researchgate.net/project/PAVICS>

Summary

- Deployment
 - Nginx + Gunicorn provide the infrastructure for scalable services
 - Birdhouse supports automatic deployment using Conda and Buildout
- Toolbox
 - Web portal and command-line tool for testing and demo of WPS services
 - Security middleware to protect the execution of WPS processes
- Get your hands dirty
 - Birdhouse Workshop: <http://birdhouse-workshop.readthedocs.io/en/latest/index.html>
 - PyWPS Workshop: <https://github.com/PyWPS/pywps-workshop>

The End

`http://bird-house.github.io/`

Appendix

The OGC Web Processing Service

- OGC open web standard for remote geo-spatial processing.
- Other widely used OGC web standards: **WMS**, **WFS**, **WCS**.
- Three basic requests:
 - *GetCapabilities*
 - *DescribeProcess*
 - *Execute*
- Three basic input/output classes:
 - *Literal*
 - *Complex* - for geo-spatial data and services
 - *BoundingBox* - for geo-spatial data extent

What does WPS provide?

- web access to your algorithms (GET request with key-value, POST request with xml)
- WPS knows about the inputs and output of a process
- processes are self-describing (GetCapabilities, DescribeProcess)
- sync and async calls (async calls with status document)
- its a standard interface ... several implementations are available (PyWPS, GeoServer, 52 North, COWS, ...)
- process definition is easy to write
- not restricted to a specific programming language
- can be used internally to provide enhanced functionality to web portals

What is PyWPS good for?

- Make your models available to the world
- Enables remote processing of complex and/or lengthy models
- Guarantees model inputs fit basic requirements (e.g. type, number)
- Guarantees interoperability of model inputs and outputs
 - using the OGC data standards
 - formalizing input/output data types

The role of a WSGI server

- WSGI - Common interface to multiple web applications frameworks
- WSGI Server - basic functions:
 - accepts HTTP requests
 - replies to HTTP requests
- WSGI provides *concurrency*, allowing multiple:
 - threads
 - workers
 - processes

Gunicorn (Green Unicorn)

- It is one of many WSGI servers out there
- Easy to configure and use with Python
- Promotes the concepts of “workers”
 - essentially OS processes
- Each worker can run on a different CPU core
 - a worker can be a Flask application instance

<http://gunicorn.org>

- Essentially a web (HTTP) server
- But more used for reverse-proxy
- Acts as a single entrance point to all requests
- Redirects requests to Gunicorn
- Can redirect to multiple Gunicorns
- Gateway to multiple servers and applications from a single URL

<http://nginx.org/>

- Python based build system
- creates application with multiple components including configuration files
- works also for non-Python parts
- using a buildout configuration
- can be extended with recipes

- An OS level virtualisation engine
- Docker runs software *containers*
 - a very light weight virtual machine
- Uses Linux kernel namespaces to isolate available resources:
 - operating environment
 - process and user IDs
 - process trees
 - network
 - mounted file systems
- Virtualisation provided by the OS itself

<https://docker.com>

Twitcher Security Proxy

